



BITS Pilani
Pilani | Dubai | Goa | Hyderabad

AIMLC ZG512 - Deep Reinforcement Learning

Session #14: Model Based Algorithms

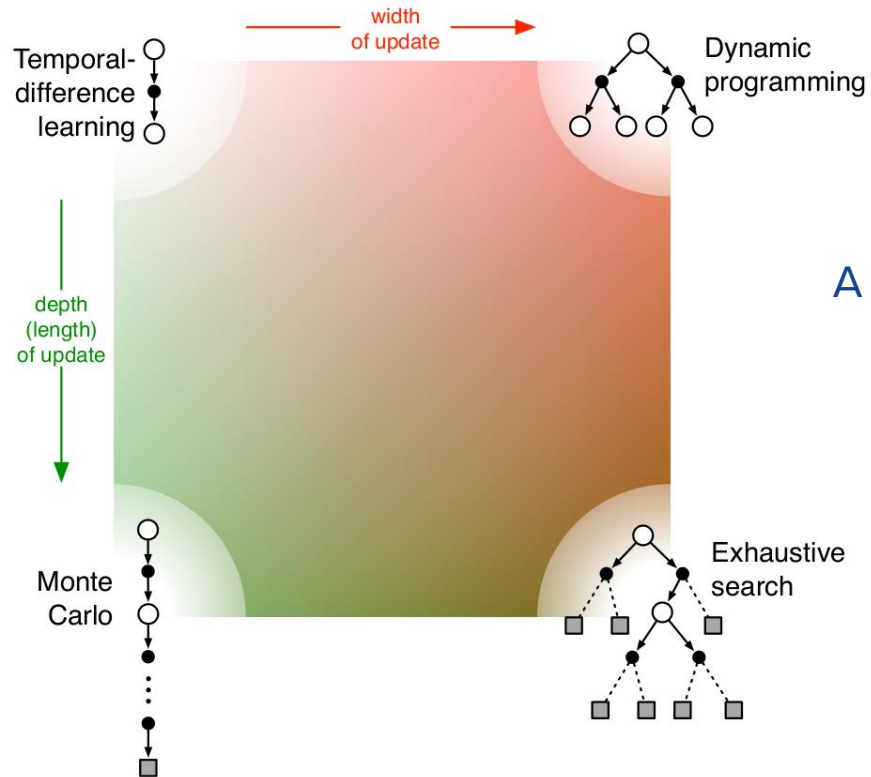


Agenda for the class

- Introduction
- Upper-Confidence-Bound [UCB] Action Selection
- Monte-Carlo Tree Search [MCTS]
- [AlphaGo](#) & [AlphaGo Zero](#) [Next Class]
- MuZero, PlaNet [Next Class]



Monte-Carlo Tree Search (MCTS)



A summary of pre-mid sem coverage !!!



Monte-Carlo Tree Search (MCTS)

Rollout Algorithms:

- Decision-time planning algorithms
- Produce Monte-Carlo estimates of action values only for each current state and for a given policy (**Rollout policy**)
- Simple, as there is no need to approximate a function over either the
 - entire state space (or)
 - state-action space
- How & Why?
 - Averaging the returns of the simulated trajectories produces estimates of $q\pi(s, a')$ for each action $a' \in A(s)$.
 - The policy selects an action in s that maximizes these estimates & then follows π
- Aim of a rollout algorithm is to improve upon the rollout policy
 - Rollout policy could be completely random !!!



Monte-Carlo Tree Search (MCTS)

Rollout Algorithms:

- Decision-time planning algorithms
- Produce Monte-Carlo estimates of action values only for each current state and for a given policy (Rollout policy)
- Simple, as there is no need to approximate a function over either the
 - entire state space (or)
 - state-action space
- How & Why?
 - Averaging the returns of the simulated trajectories produces estimates of $q\pi(s, a')$ for each action $a' \in A(s)$.
 - The policy selects an action in s that maximizes these estimates & then follows π
- Aim of a rollout algorithm is to improve upon the rollout policy
 - Rollout policy could be completely random !!!
- MCTS is a recent and strikingly successful example of decision-time planning
- An enhanced rollout algorithm
 - Accumulates value estimates obtained from the simulations to successively direct simulations toward more highly-rewarding trajectories



Monte-Carlo Tree Search (MCTS)


How MCTS works?

- MCTS is *executed* after encountering each new state (s)
 - [?] to select the agent's action for s
- *Each execution is an iterative process* that simulates many trajectories starting from s and
 - running to a terminal state (or)
 - until discounting makes any further reward negligible to the return
- Focus on multiple simulations starting at s by extending the initial portions of trajectories that have received high evaluations from earlier simulations.




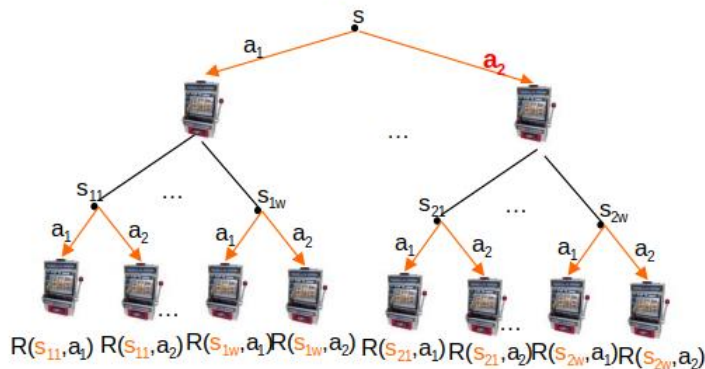
Monte-Carlo Tree Search (MCTS)

Real world state/action sequence



Build look-ahead tree

Build look-ahead tree




.....

.....



Monte-Carlo Tree Search (MCTS)

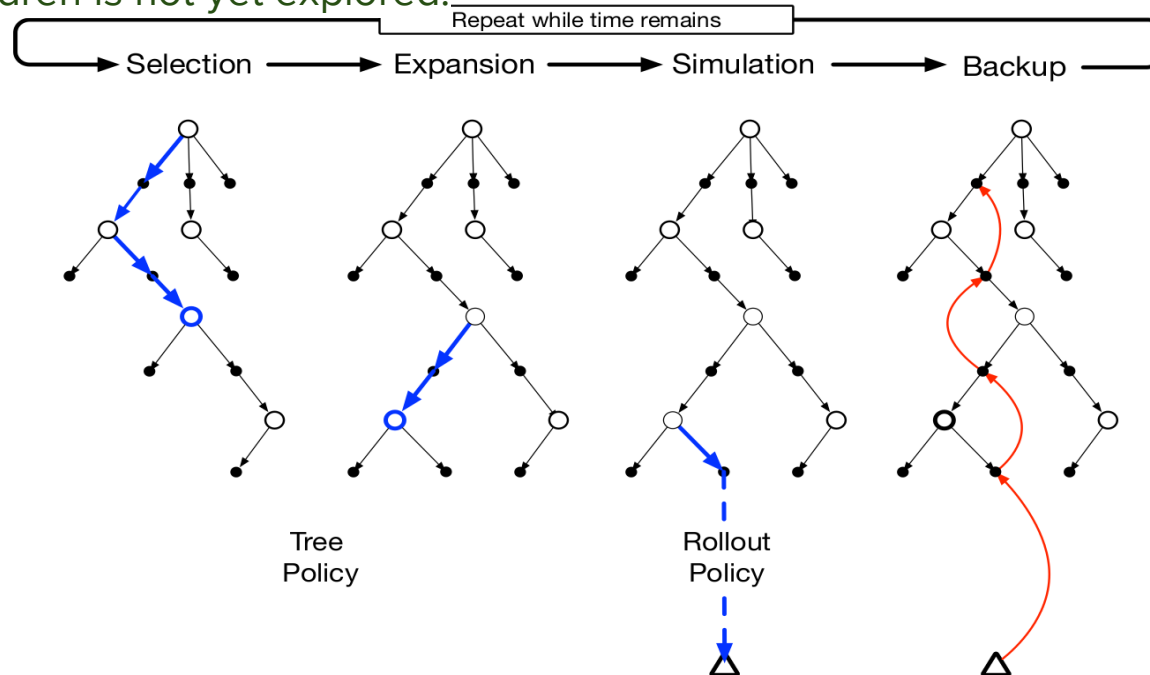


$$S_i = x_i + C \sqrt{\frac{\ln(t)}{n_i}}$$



Monte-Carlo Tree Search (MCTS) -- Selection

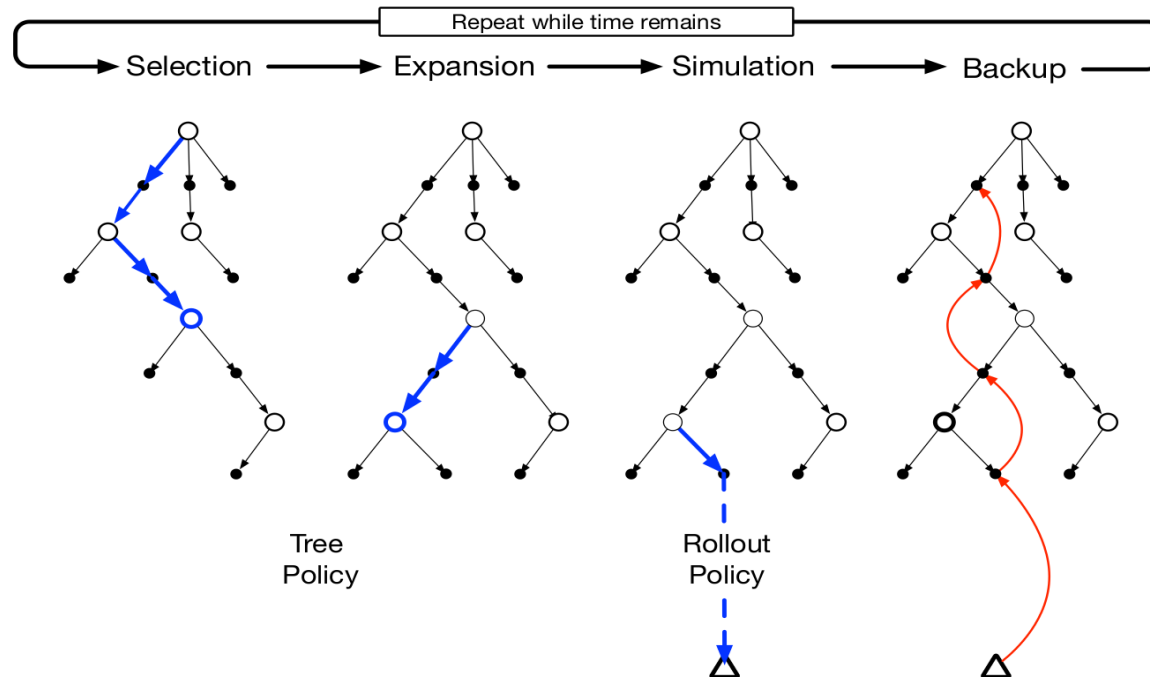
Select: Select a single node in the tree that is *not fully expanded*. By this, we mean at least one of its children is not yet explored.





Monte-Carlo Tree Search (MCTS) -- Expansion


Expand: Expand this node by applying one available action (as defined by the MDP) from the node.





Monte-Carlo Tree Search (MCTS) -- Simulation

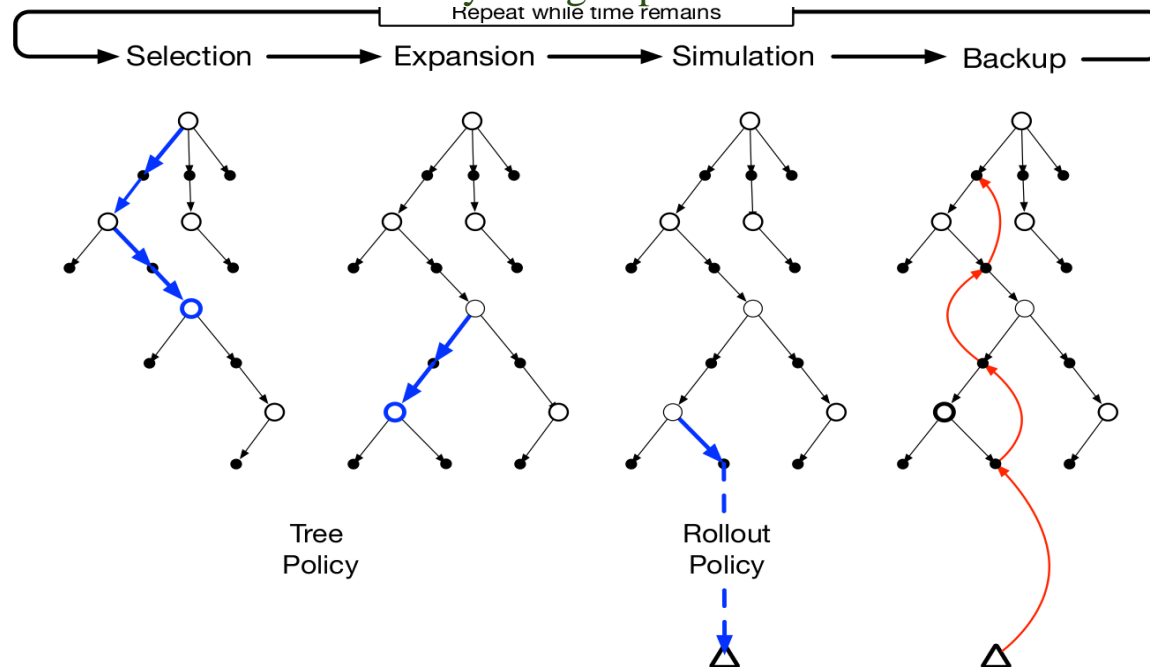
Simulation: From one of the outcomes of the expanded, perform a complete random simulation to a terminating state.





Monte-Carlo Tree Search (MCTS) -- Backup

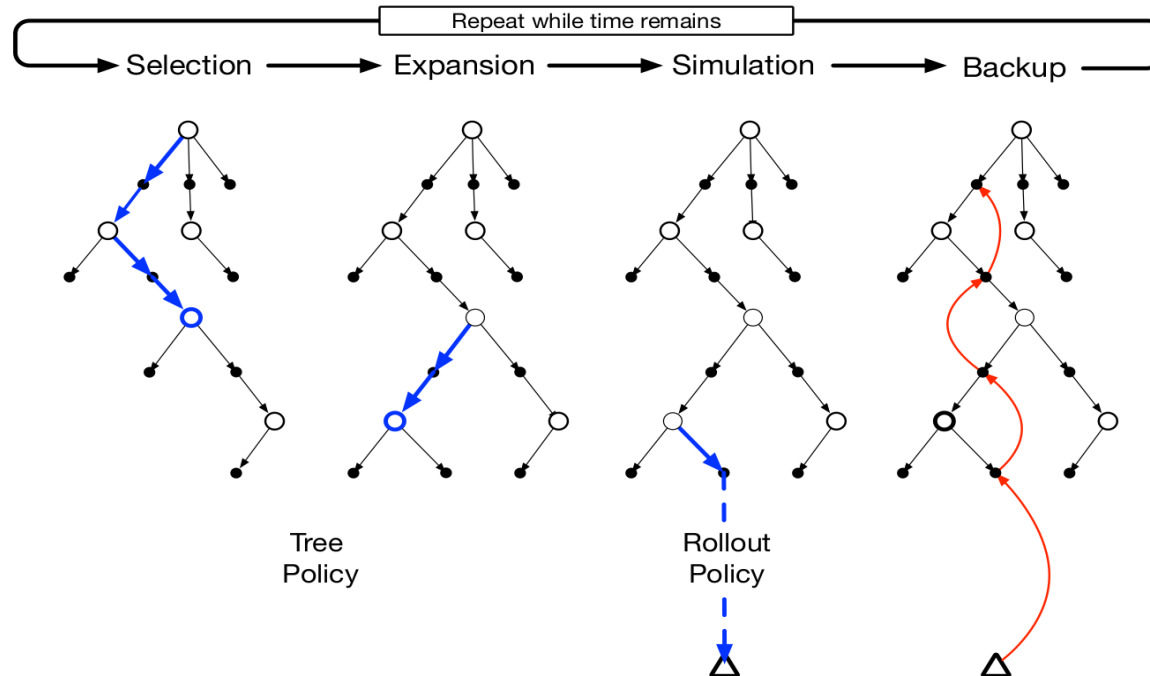
Backup/ Backpropagate: The value of the node is *back propagated* to the root node, updating the value of each ancestor node on the way using expected value






Monte-Carlo Tree Search (MCTS) -- Summarizing

Comments on the overall approach,,,







Monte-Carlo Tree Search (MCTS) -- Selection



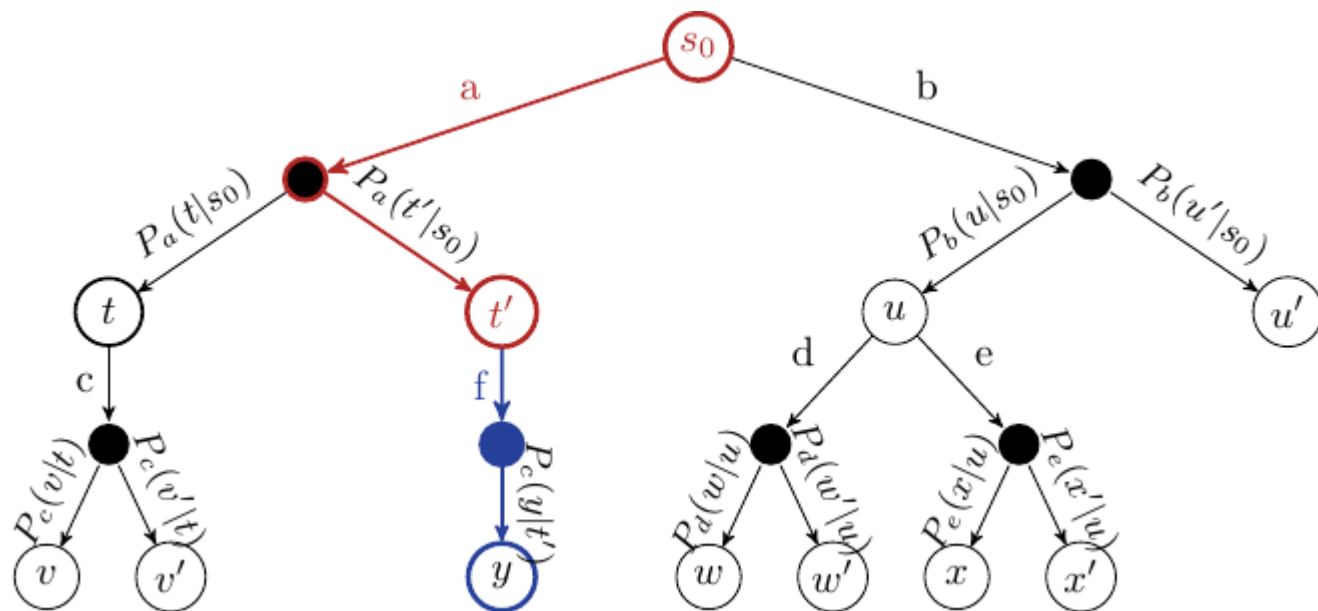


Monte-Carlo Tree Search (MCTS) -- Selection



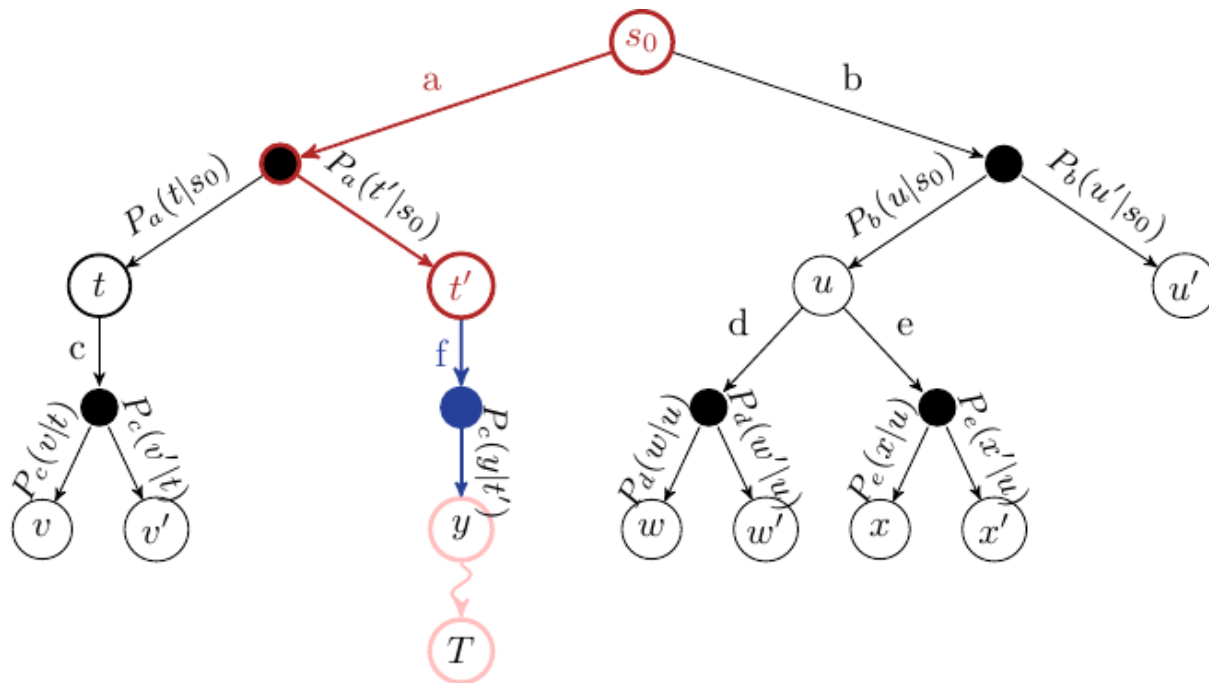


Monte-Carlo Tree Search (MCTS) -- Expansion



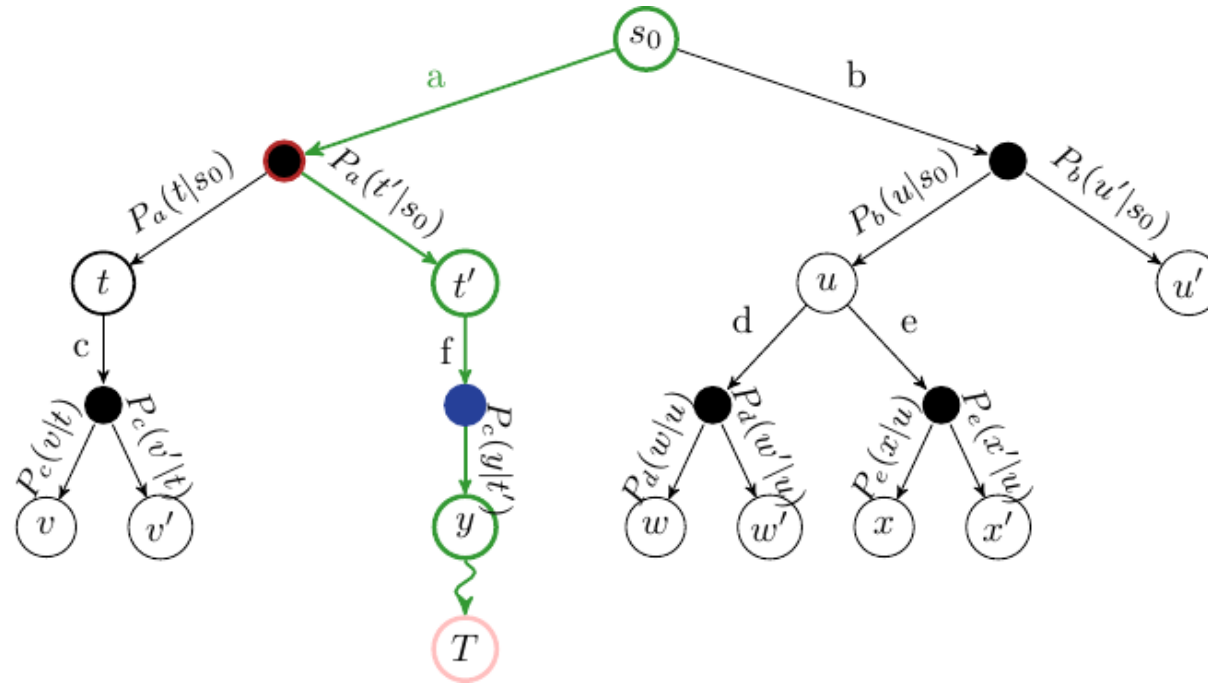


Monte-Carlo Tree Search (MCTS) -- Simulation





Monte-Carlo Tree Search (MCTS) -- Backup





Monte-Carlo Tree Search (MCTS)

Algorithm – Monte-Carlo Tree Search

Input: MDP $M = \langle S, s_0, A, P_a(s' | s), r(s, a, s') \rangle$, base value function Q , time limit T .

Output: updated Q-function Q

while $currentTime < T$

$selected_node \leftarrow \text{Select}(s_0)$

$child \leftarrow \text{Expand}(selected_node)$ – expand and choose a child to simulate

$G \leftarrow \text{Simulate}(child)$ – simulate from $child$

$\text{Backpropagate}(selected_node, child, G)$

return Q



Monte-Carlo Tree Search (MCTS)

Function – $\text{Select}(s : S)$

Input: state s

Output: unexpanded state

while s is fully expanded

 Select action a to apply in s using a multi-armed bandit algorithm

 Choose one outcome s' according to $P_a(s' \mid s)$

$s \leftarrow s'$

return s



Monte-Carlo Tree Search (MCTS)

Function – $\text{Expand}(s : S)$

Input: state s

Output: expanded state s'

Select an action a from s to apply

Expand one outcome s' according to the distribution $P_a(s' | s)$ and observe reward r

return s'



Monte-Carlo Tree Search (MCTS)

🔔 Procedure – Backpropagation($s : S; a : A$)

Input: state-action pair (s, a)

Output: none

do

$$N(s, a) \leftarrow N(s, a) + 1$$

$$G \leftarrow r + \gamma G$$

$$Q(s, a) \leftarrow Q(s, a) + \frac{1}{N(s, a)} [G - Q(s, a)]$$


$$s \leftarrow \text{parent of } s$$

$$a \leftarrow \text{parent action of } s$$

while $s \neq s_0$



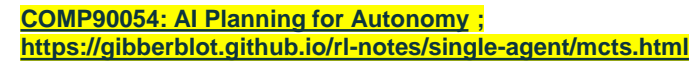
Monte-Carlo Tree Search (MCTS)



Before backpropagation

$$Q(s, a) = 18$$

$$Q(t, f) = 0$$


$$\begin{array}{rcl} Q(s, a) & = & 18 \\ Q(t, f) & = & 0 \end{array}$$
$$\begin{aligned} Q(y, g) &= \gamma^2 \times 31.25 \quad (\text{simulation is 3 steps long and receives reward of 31.25}) \\ &= 20 \\ N(t, f) &\leftarrow N(t, f) + 1 = N(y) + N(y') + N(y'') + 1 = 2 \\ Q(t, f) &= Q(t, f) + \frac{1}{N(t, f)}[r + \gamma G - Q(t, f)] \\ &= 0 + \frac{1}{2}[0 + 0.8 \cdot 20 - 0] \\ &= 8 \\ N(s, a) &\leftarrow N(s, a) + 1 = N(t) + N(t') + 1 = 5 \\ Q(s, a) &= Q(s, a) + \frac{1}{N(s, a)}[r + \gamma G - Q(s, a)] \\ &= 18 + \frac{1}{5}[6 + 0.8 \cdot (0.8 \cdot 20) - 18] \\ &= 18 + \frac{1}{5}[6 + 12.8 - 18] \\ &= 18.16 \end{aligned}$$





Upper-Confidence-Bound Action Selection

- ϵ -greedy action selection forces the non-greedy actions to be tried,
Indiscriminately, with no preference for those that are nearly greedy or particularly uncertain
- It would be better to select among the non-greedy actions according to their potential for actually being optimal
Take into account both how close their estimates are to being maximal and the uncertainties in those estimates.

$$A_t \doteq \arg \max_a \left[Q_t(a) + c \sqrt{\frac{\ln t}{N_t(a)}} \right]$$



Upper-Confidence-Bound Action Selection

- Each time a is selected the uncertainty is presumably reduced
- Each time an action other than a is selected, t increases but $N_t(a)$ does not; because t appears in the numerator, the uncertainty estimate increases.
- Actions with lower value estimates, or that have already been selected frequently, will be selected with decreasing frequency over time

Action Value at time t for a

Confidence Level

Measure of Uncertainty

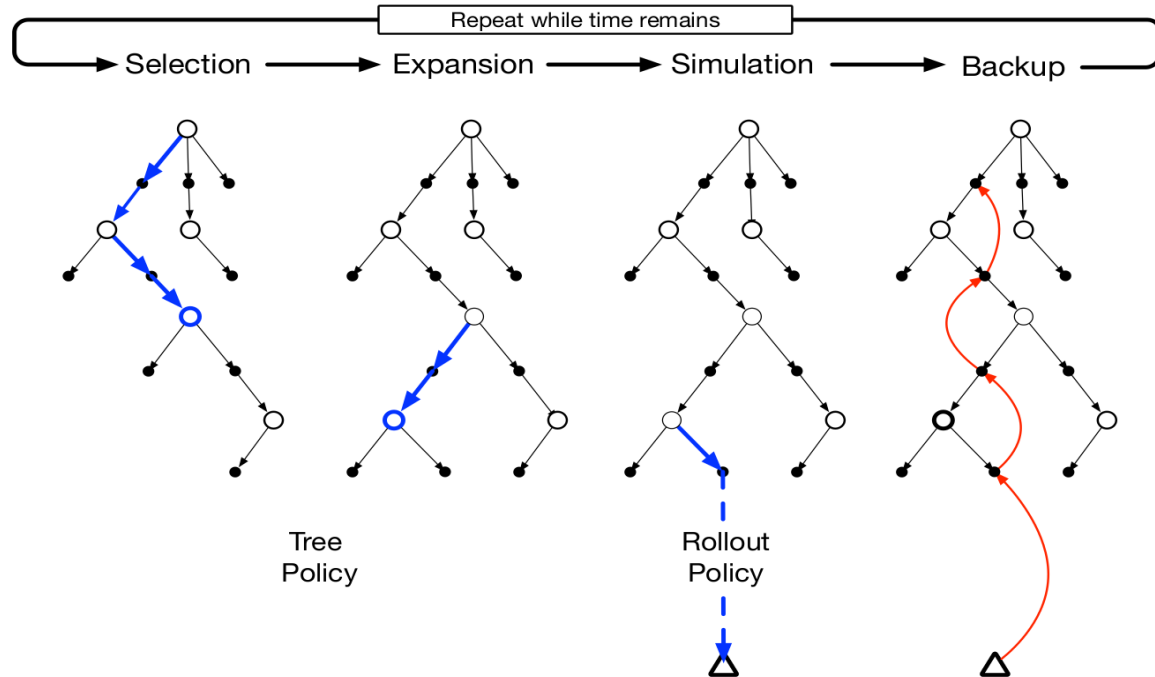
$$A_t \doteq \arg \max_a \left[Q_t(a) + c \sqrt{\frac{\ln t}{N_t(a)}} \right]$$



Monte-Carlo Tree Search (MCTS)

Can the selection of action in Tree policy use UCB?

$$S_i = x_i + C \sqrt{\frac{\ln(t)}{n_i}}$$






Monte-Carlo Tree Search (MCTS)

Can the selection of action
in Tree policy use UCB?

Upper Confidence Trees (UCT):
MCTS with UCB for Tree policy





Required Readings and references

1. <https://rl-lab.com/#play>
2. <https://www.aionlinecourse.com/tutorial/machine-learning/upper-confidence-bound-%28ucb%29>
3. <https://towardsdatascience.com/monte-carlo-tree-search-in-reinforcement-learning-b97d3e743d0f>
4. <https://gibberblot.github.io/rl-notes/single-agent/mcts.html>
5. <https://towardsdatascience.com/alphazero-chess-how-it-works-what-sets-it-apart-and-what-it-can-tell-us-4ab3d2d08867>
6. <https://medium.com/geekculture/muzero-explained-a04cb1bad4d4>
7. <https://towardsdatascience.com/everything-you-need-to-know-about-googles-new-planet-reinforcement-learning-network-144c2ca3f284>
8. <https://blog.research.google/2019/02/introducing-planet-deep-planning.html?m=1>



BITS Pilani
Pilani | Dubai | Goa | Hyderabad

Thank you