



**BITS Pilani**  
Pilani | Dubai | Goa | Hyderabad

# Deep Reinforcement Learning

2022-23 Second Semester, M.Tech (AIML)

## Session #15: Imitation Learning

### Instructors :

1. Prof. S. P. Vimal ([vimalsp@wilp.bits-pilani.ac.in](mailto:vimalsp@wilp.bits-pilani.ac.in)),
2. Prof. Sangeetha Viswanathan ([sangeetha.viswanathan@pilani.bits-pilani.ac.in](mailto:sangeetha.viswanathan@pilani.bits-pilani.ac.in))



# Agenda for the classes

- Imitation Learning
  - Behaviour Cloning
  - Inverse RL




# Imitation Learning in a Nutshell


**Given:** demonstrations or demonstrator

**Goal:** train a policy to mimic demonstrations


Expert Demonstrations



State/Action Pairs




Learning






# Ingredients of Imitation Learning



Demonstrations or Demonstrator



Environment / Simulator



Policy Class



Loss Function



Learning Algorithm



# Some Interesting Examples

- ALVINN

<https://www.ri.cmu.edu/publications/alvin-an-autonomous-land-vehicle-in-a-neural-network/>

Dean Pomerleau et al., 1989-1999 <https://www.youtube.com/watch?v=iIP4aPDTBPE>

- Helicopter Acrobatics

**Learning for Control from Multiple Demonstrations** - Adam Coates, Pieter Abbeel, Andrew Ng, ICML 2008

**An Application of Reinforcement Learning to Aerobatic Helicopter Flight** - Pieter Abbeel, Adam Coates, Morgan Quigley, Andrew Y. Ng, NIPS 2006

<https://www.youtube.com/watch?v=0JL04JJjocc>

- Ghosting ( Sports Analytics) - Next Slide.



## Some Interesting Examples





# Some Interesting Examples

## What's Hidden in the Hidden Layers?

*The contents can be easy to find with a geometrical problem,  
but the hidden layers have yet to give up all their secrets*

David S. Touretzky and Dean A. Pomerleau

AUGUST 1989 • B Y T E 231

tions, we fed the network road images taken under a wide variety of viewing angles and lighting conditions. It would be impractical to try to collect thousands of real road images for such a data set. Instead, we developed a synthetic road-image generator that can create as many training examples as we need.

To train the network, 1200 simulated road images are presented 40 times each, while the weights are adjusted using the back-propagation learning algorithm. This takes about 30 minutes on Carnegie Mellon's Warp systolic-array supercomputer. (This machine was designed at Carnegie Mellon and is built by General Electric. It has a peak rate of 100 million floating-point operations per second and can compute weight adjustments for back-propagation networks at a rate of 20 million connections per second.)

Once it is trained, ALVINN can accurately drive the NAVLAB vehicle at about 3½ miles per hour along a path through a wooded area adjoining the Carnegie Mellon campus, under a variety of weather and lighting conditions. This speed is nearly twice as fast as that achieved by non-neural-network algorithms running on the same vehicle. Part of the reason for this is that the forward pass of a back-propagation network can be computed quickly. It takes about 200

milliseconds on the Sun-3/160 workstation installed on the NAVLAB.

The hidden-layer representations ALVINN develops are interesting. When trained on roads of a fixed width, the net-

work chooses a representation in which hidden units act as detectors for complete roads at various positions and orientations. When trained on roads of variable

*continued*



Photo 1: The NAVLAB autonomous navigation test-bed vehicle and the road used for trial runs.



# Some Interesting Examples



**Learning for Control from Multiple Demonstrations** - Adam Coates, Pieter Abbeel, Andrew Ng, ICML 2008

**An Application of Reinforcement Learning to Aerobatic Helicopter Flight** - Pieter Abbeel, Adam Coates, Morgan Quigley, Andrew Y. Ng, NIPS 2006





# Ghosting

## Data Driven Ghosting using Deep Imitation Learning

Hoang M. Le et al., SSAC 2017



English Premier League  
2012-2013

Match date: 04/05/2013

<https://www.youtube.com/watch?v=WI-WL2cj0CA>



# Some Interesting Examples





# Notation & Set-up

State:  $s$  (sometimes  $x$ ) (\*\*state may only be partially observed)

Action:  $a$  (sometimes  $y$ )

Policy:  $\pi_{\theta}$  (sometimes  $h$ )

- Policy maps states to actions:  $\pi_{\theta}(s) \rightarrow a$
- ...or distributions over actions:  $\pi_{\theta}(s) \rightarrow P(a)$

State Dynamics:  $P(s'|s,a)$

- Typically not known to policy.
- Essentially the simulator/environment



# Notation & Set-up

Rollout: sequentially execute  $\pi(s_0)$  on an initial state

- Produce trajectory  $\tau=(s_0,a_0,s_1,a_1,\dots)$

$P(\tau|\pi)$ : distribution of trajectories induced by a policy

1. Sample  $s_0$  from  $P_0$  (distribution over initial states), initialize  $t = 1$ .
2. Sample action  $a_t$  from  $\pi(s_{t-1})$
3. Sample next state  $s_t$  from applying  $a_t$  to  $s_{t-1}$  (requires access to environment)
4. Repeat from Step 2 with  $t=t+1$

$P(s|\pi)$ : distribution of states induced by a policy

- Let  $P_t(s|\pi)$  denote distribution over  $t$ -th state
- $P(s|\pi) = (1/T)\sum_t P_t(s|\pi)$



# Example #1: Racing Game

(Super Tux Kart)



$s$  = game screen

$a$  = turning angle

Training set:  $D = \{r := (s, a)\}$  from  $\pi^*$

- $s$  = sequence of  $s$
- $a$  = sequence of  $a$

**Goal:** learn  $\pi_\theta(s) \rightarrow a$





## Example #2: Basketball Trajectories


$\mathbf{s}$  = location of players & ball

$\mathbf{a}$  = next location of player

Training set:  $D=\{r:=\langle \mathbf{s}, \mathbf{a} \rangle\}$  from  $\pi^*$

- $\mathbf{s}$  = sequence of  $\mathbf{s}$
- $\mathbf{a}$  = sequence of  $\mathbf{a}$

**Goal:** learn  $\pi_\theta(\mathbf{s}) \rightarrow \mathbf{a}$





**Behavioral Cloning** = Reduction to Supervised Learning (Ignoring regularization for brevity.)

**Write:**

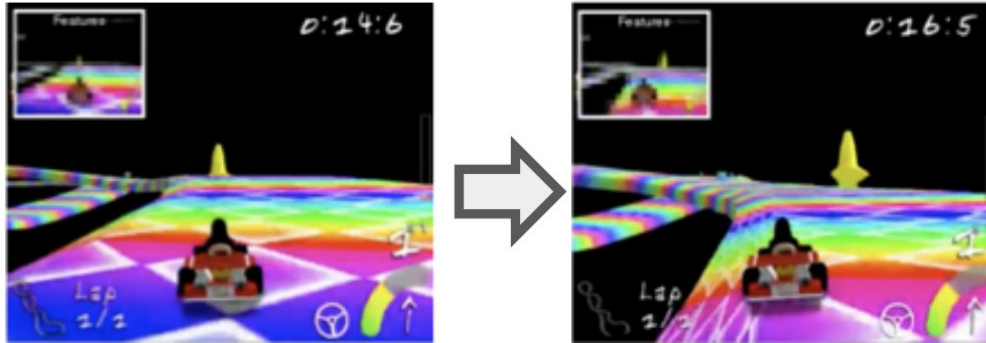


# Behavioral Cloning vs. Imitation Learning





# Limitations of Behavioral Cloning



$\pi_{\theta}$  makes a mistake

New state sampled not from  $P^*$ !

**Worst case is catastrophic!**

Images from Stephane Ross

IID Assumption  
(Supervised Learning)


$P^*$



$(s, a^*)$


Reality

$P_0$





# Limitations of Behavioral Cloning



Compounding Errors

Data distribution mismatch!

In supervised learning,  $(x, y) \sim D$  during train **and** test. In MDPs:

- Train:  $s_t \sim D_{\pi^*}$
- Test:  $s_t \sim D_{\pi_\theta}$



# When to use Behavioral Cloning?

## Advantages

- Simple
- Simple
- Efficient

## Use When:

- 1-step deviations not too bad
- Learning reactive behaviors
- Expert trajectories “cover” state space

## Disadvantages

- Distribution mismatch between training and testing
- No long term planning

## Don't Use When:

- 1-step deviations can lead to catastrophic error
- Optimizing long-term objective (at least not without a stronger model)



# Types of Imitation Learning

## Behavioral Cloning

$$\operatorname{argmin}_{\theta} E_{(s,a^*) \sim P^*} L(a^*, \pi_{\theta}(s))$$

**Works well when  $P^*$  close to  $P_{\theta}$**

## Inverse RL


Learn  $r$  such that:

$$\pi^* = \operatorname{argmax}_{\theta} E_{s \sim P(s|\theta)} r(s, \pi_{\theta}(s))$$

**RL problem**

**Assumes learning  $r$  is statistically easier than directly learning  $\pi^*$**

## Direct Policy Learning via Interactive Demonstrator



**Requires Interactive Demonstrator  
(BC is 1-step special case)**



# Interactive Expert

Can query expert at any state


Construct loss function

- $L(\pi^*(s), \pi(s))$

Typically applied to rollout trajectories

- $s \sim P(s|\pi)$

Driving example:  $L(\pi^*(s), \pi(s)) = (\pi^*(s) - \pi(s))^2$



Example from Super Tux Kart  
(Image courtesy of Stephane Ross)

Expert provides feedback on  
state visited by policy



# DAGGER: Dataset Aggregation

Initialize  $\mathcal{D} \leftarrow \emptyset$ .

Initialize  $\hat{\pi}_1$  to any policy in  $\Pi$ .

**for**  $i = 1$  **to**  $N$  **do**

Let  $\pi_i = \beta_i \pi^* + (1 - \beta_i) \hat{\pi}_i$ .

Sample  $T$ -step trajectories using  $\pi_i$ .

Get dataset  $\mathcal{D}_i = \{(s, \pi^*(s))\}$  of visited states by  $\pi_i$   
and actions given by expert.

Aggregate datasets:  $\mathcal{D} \leftarrow \mathcal{D} \cup \mathcal{D}_i$ .

Train classifier  $\hat{\pi}_{i+1}$  on  $\mathcal{D}$ .

**end for**

**Return** best  $\hat{\pi}_i$  on validation.

$\beta_i$ : a decreasing coefficient s.t.  
 $\frac{1}{N} \sum_{i=1}^N \beta_i \rightarrow 0$  as  $N \rightarrow \infty$

- Idea: Get more labels of the expert action along the path taken by the policy computed by behavior cloning
- Obtains a stationary deterministic policy with good performance under its induced state distribution



# Inverse RL

- What if we don't have an online demonstrator?
  - We only have access to an offline set of demonstrated trajectories
- Behavioral cloning is not robust
  - Suffers from overfitting
  - We know what to do in observed states but can't generalize well to other states
- How can we learn to mimic the demonstrator in a general way?
  - Learn the demonstrator's objective (reward) function
  - Apply RL




# Inverse RL

- What if we don't have an online demonstrator?
  - We only have access to an offline set of demonstrated trajectories
- Behavioral cloning is not robust
  - Suffers from overfitting
  - We know what to do in observed states but can't generalize well to other states
- How can we learn to mimic the demonstrator in a general way?
  - Learn the demonstrator's objective (reward) function
  - Apply RL






# Inverse RL



Muybridge (c. 1870)



Mombaur et al. '09



Li & Todorov '06



Ziebart '08

$$\mathbf{a}_1, \dots, \mathbf{a}_T = \arg \max_{\mathbf{a}_1, \dots, \mathbf{a}_T} \sum_{t=1}^T r(\mathbf{s}_t, \mathbf{a}_t)$$

$$\mathbf{s}_{t+1} = f(\mathbf{s}_t, \mathbf{a}_t)$$


$$\pi = \arg \max_{\pi} E_{\mathbf{s}_{t+1} \sim p(\mathbf{s}_{t+1} | \mathbf{s}_t, \mathbf{a}_t), \mathbf{a}_t \sim \pi(\mathbf{a}_t | \mathbf{s}_t)} [r(\mathbf{s}_t, \mathbf{a}_t)]$$

$$\mathbf{a}_t \sim \pi(\mathbf{a}_t | \mathbf{s}_t)$$

optimize this to explain the data




# Inverse RL



Muybridge (c. 1870)



Mombaur et al. '09



Li & Todorov '06



Ziebart '08

$$\mathbf{a}_1, \dots, \mathbf{a}_T = \arg \max_{\mathbf{a}_1, \dots, \mathbf{a}_T} \sum_{t=1}^T r(\mathbf{s}_t, \mathbf{a}_t)$$

$$\mathbf{s}_{t+1} = f(\mathbf{s}_t, \mathbf{a}_t)$$

$$\pi = \arg \max_{\pi} E_{\mathbf{s}_{t+1} \sim p(\mathbf{s}_{t+1} | \mathbf{s}_t, \mathbf{a}_t), \mathbf{a}_t \sim \pi(\mathbf{a}_t | \mathbf{s}_t)} [r(\mathbf{s}_t, \mathbf{a}_t)]$$

$$\mathbf{a}_t \sim \pi(\mathbf{a}_t | \mathbf{s}_t)$$

optimize this to explain the data




# Inverse RL

## The imitation learning perspective

Standard imitation learning:

- copy the *actions* performed by the expert
- no reasoning about outcomes of actions



Human imitation learning:



- copy the *intent* of the expert
- might take very different actions!





# Inverse RL

The reinforcement learning perspective




**what is the reward?**



# Inverse RL





Infer reward functions from **demonstrations**



➡  $r(s, a)$

by itself, this is an **underspecified** problem

many reward functions can explain the **same** behavior





# Inverse RL

"forward" reinforcement learning

given:

states  $\mathbf{s} \in \mathcal{S}$ , actions  $\mathbf{a} \in \mathcal{A}$

(sometimes) transitions  $p(\mathbf{s}'|\mathbf{s}, \mathbf{a})$

reward function  $r(\mathbf{s}, \mathbf{a})$

learn  $\pi^*(\mathbf{a}|\mathbf{s})$

inverse reinforcement learning

given:

states  $\mathbf{s} \in \mathcal{S}$ , actions  $\mathbf{a} \in \mathcal{A}$

(sometimes) transitions  $p(\mathbf{s}'|\mathbf{s}, \mathbf{a})$

samples  $\{\tau_i\}$  sampled from  $\pi^*(\tau)$

learn  $r_\psi(\mathbf{s}, \mathbf{a})$

reward parameters


...and then use it to learn  $\pi^*(\mathbf{a}|\mathbf{s})$

neural net reward function:

linear reward function:

$$r_\psi(\mathbf{s}, \mathbf{a}) = \sum_i \psi_i f_i(\mathbf{s}, \mathbf{a}) = \psi^T \mathbf{f}(\mathbf{s}, \mathbf{a})$$

$\mathbf{s}$   
 $\mathbf{a}$



$r_\psi(\mathbf{s}, \mathbf{a})$

parameters  $\psi$


# GAN




Zhu et al. '17



Arjovsky et al. '17



Isola et al. '17




samples from  $p_{\theta}(\mathbf{x})$

data (“demonstrations”)



$D(\mathbf{x}) = p_{\psi}(\text{real image}|\mathbf{x})$

samples from  $p^*(\mathbf{x})$

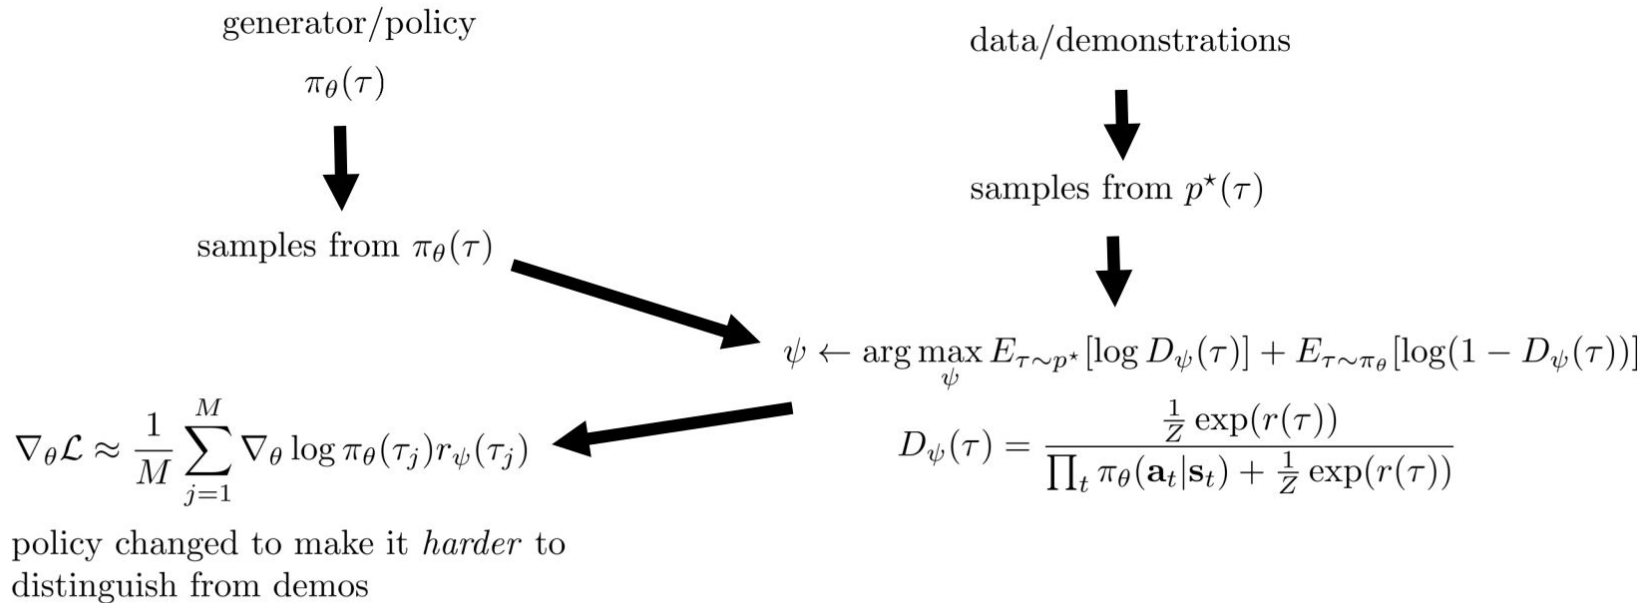


$$\psi = \arg \max_{\psi} \frac{1}{N} \sum_{\mathbf{x} \sim p^*} \log D_{\psi}(\mathbf{x}) + \frac{1}{M} \sum_{\mathbf{x} \sim p_{\theta}} \log(1 - D_{\psi}(\mathbf{x}))$$

$$\theta \leftarrow \arg \max_{\theta} E_{\mathbf{x} \sim p_{\theta}} \log D_{\psi}(\mathbf{x})$$



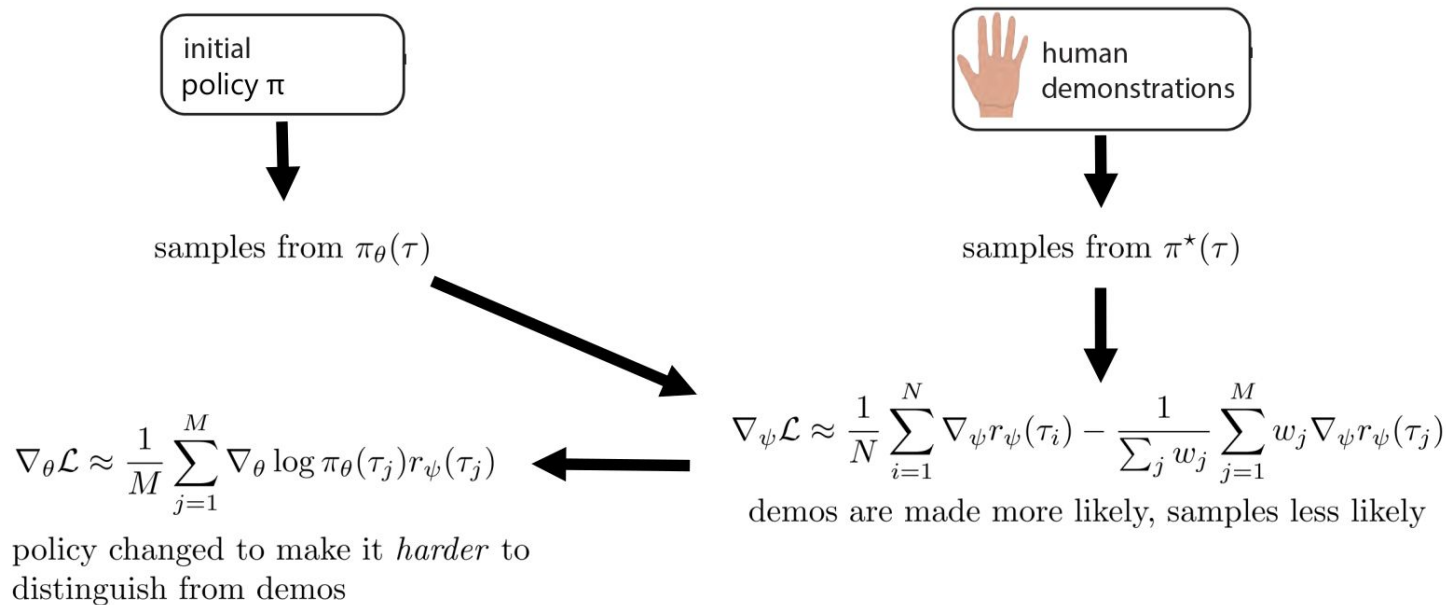
# Inverse RL as GAN







# Inverse RL as GAN





## Required Readings and references

1. [Human-in-the-Loop Deep Reinforcement Learning with Application to Autonomous Driving.](#) Jingda Wu, Zhiyu Huang, Chao Huang, Zhongxu Hu, Peng Hang, Yang Xing, Chen Lv\*



**BITS Pilani**  
Pilani | Dubai | Goa | Hyderabad

Thank you